

Lending Club Loan Data Analysis

PG AI – Deep Learning with Tensorflow and Keras Project:

Project1: Lending Club Loan Data Analysis

Writeup

We have the subject of building a deep learning model to predict the chance of default for future loans.

The model will be based on a dataset which presents some historical data on clients such as The number of days the borrower has had a credit line, the borrower's number of inquiries by creditors in the last 6 months and other information.

Before beginning the construction of the deep learning model, we first carry out the EAD and feature engineering operations. In the next step we look at the performance of some machine learning algorithms to consider them as a basis for reflection on our deep learning model to be built. Before attacking our objective, we will try to reduce the features by referring to `feature_importances_` of `RandomForestClassifier`.

To build our deep learning model, we will use two overfitting techniques **Dropout** **Regularization** and **L2 Regularization**. For more efficiency, the operation which consists of to lower the learning rate as the training progresses is ensured by **the inverseTimeDecay** function. And to finalize our model we will use a **Callback** function with a target.

the Python libraries that will be useful to us can be found at the end of this document with this [link](#), and after loading the dataset :

```
import pandas as pd
df_init = pd.read_csv('/content/loan_data.csv')
#
df=df_init.copy(deep=True)
# Top 5 records
df.head()
```

Rachid LAMJOUN

PG AI Program

	credit.policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1	0	0	0	0
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7	0	0	0	0
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6	1	0	0	0
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2	1	0	0	0
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5	0	1	0	0

We give some necessary information to better know our dataset:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose                9578 non-null   object
2   int.rate               9578 non-null   float64
3   installment            9578 non-null   float64
4   log.annual.inc         9578 non-null   float64
5   dti                    9578 non-null   float64
6   fico                   9578 non-null   int64
7   days.with.cr.line      9578 non-null   float64
8   revol.bal              9578 non-null   int64
9   revol.util             9578 non-null   float64
10  inq.last.6mths         9578 non-null   int64
11  delinq.2yrs            9578 non-null   int64
12  pub.rec                9578 non-null   int64
13  not.fully.paid         9578 non-null   int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

We have, therefore, a dataset with 9578 entries and without null value and 14 columns. Now let's analyze the type of these columns :

```
df.dtypes.value_counts()
```

```
int64      7
float64    6
object     1
dtype: int64
```

```
df.select_dtypes('object').head()
```

purpose	
0	debt_consolidation
1	credit_card
2	debt_consolidation
3	debt_consolidation
4	credit_card

A single column of type 'object':

```
df['purpose'].value_counts()
```

```
debt_consolidation    3957  
all_other             2331  
credit_card           1262  
home_improvement      629  
small_business         619  
major_purchase        437  
educational           343  
Name: purpose, dtype: int64
```

For the following we will use the function [get_value_counts\(\)](#) which gives for a given type the name of values that appear in the columns of this type.

```
# type= object  
lst_t=get_value_counts(df,'object')  
lst_t
```

```
[ ('purpose', 7) ]
```

We will then have to do the transformation categorical values into numerical values of this column. Before we analyze the other types of columns :

```
# int64  
df.select_dtypes('int64').head()
```

	credit.policy	fico	revol.bal	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
0	1	737	28854	0	0	0	0
1	1	707	33623	0	0	0	0
2	1	682	3511	1	0	0	0
3	1	712	33667	1	0	0	0
4	1	667	4740	0	1	0	0

```
# type= int64
lst_t=get_value_counts(df,'int64')
lst_t
```

```
[('credit.policy', 2),
 ('fico', 44),
 ('revol.bal', 7869),
 ('inq.last.6mths', 28),
 ('delinq.2yrs', 11),
 ('pub.rec', 6),
 ('not.fully.paid', 2)]
```

```
# type= float64
lst_t=get_value_counts(df,'float64')
lst_t
```

```
[('int.rate', 249),
 ('installment', 4788),
 ('log.annual.inc', 1987),
 ('dti', 2529),
 ('days.with.cr.line', 2687),
 ('revol.util', 1035)]
```

In order to start studying the models let's apply the **get_dummies()** function of pandas and analyze the result :

```
df1=pd.get_dummies(df)
df1.head()
```

Rachid LAMJOUN

PG AI Program

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	...
0	1	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1	0	...
1	1	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7	0	...
2	1	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6	1	...
3	1	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2	1	...
4	1	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5	0	...

```
df1.dtypes.value_counts()
```

```
int64      7
float64    6
object     1
dtype: int64
```

```
# reminder
```

```
df.dtypes.value_counts()
```

```
int64      7
float64    6
object     1
dtype: int64
```

```
# uint8
```

```
df1.select_dtypes('uint8').columns
```

```
Index(['purpose_all_other', 'purpose_credit_card',
      'purpose_debt_consolidation', 'purpose_educational',
      'purpose_home_improvement', 'purpose_major_purchase',
      'purpose_small_business'],
      dtype='object')
```

We had the replacement of the column purpose by these columns whose name is prefixed by purpose.

```
# type= uint8
```

```
lst_t=get_value_counts(df1,'uint8')
```

```
lst_t
```

```
[('purpose_all_other', 2),  
 ('purpose_credit_card', 2),  
 ('purpose_debt_consolidation', 2),  
 ('purpose_educational', 2),  
 ('purpose_home_improvement', 2),  
 ('purpose_major_purchase', 2),  
 ('purpose_small_business', 2)]
```

Only values 0 and 1 for these columns :

```
# type= uint8  
lst_t=df1.select_dtypes('uint8').values  
ll=[]  
for i in range(len(lst_t)):  
    ll=ll+list(lst_t[i,:])  
set(ll)
```

```
{0, 1}
```

We define the X and the Y (Target) to then, define also, the Train and Validation data:

```
# def of X, Y  
Y=df1[df1.columns[0]]  
X=df1[df1.columns[1:]]  
X.shape, Y.shape, len(df1), len(df1.columns)
```

```
((9578, 19), (9578,), 9578, 20)
```


We check features dependency :

```
# Significant dependency  
corr_matrix_X = X.corr()  
print('corr_matrix_X.shape')  
print(corr_matrix_X.shape)  
#  
coef_dep=0.85  
#  
if len(verify_dependency(corr_matrix_X,coef_dep))==0:  
    print('\nNo dependency:', verify_dependency(corr_matrix_X,coef_dep))  
else:  
    print('\nDependency:....', verify_dependency(corr_matrix_X,coef_dep))  
#  
print('\n')  
import seaborn as sns
```

```
sns.heatmap(corr_matrix_X)
```

```
corr_matrix_X.shape  
(19, 19)
```

```
No dependency: []
```



At this level we split the data with 25% the validation data:

```
# Split() --> define X_train,X_test,y_train,y_test  
from sklearn.model_selection import train_test_split  
#  
in_test_size=25/100  
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=in_test_size,random_state=0)  
#  
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((7183, 19), (2395, 19), (7183,), (2395,))
```

Before the build of our deep learning model, we look at the performance of some classical machine learning models by referring to function [test_classifierModels_list\(\)](#):

```
# list classifiers  
lst_algo=['AdaBoost','SVM','KNN','DecisionTreeClassifier','RandomForestClassifier']  
#  
dico_name_model=test_classifierModels_list(lst_algo,X_train,X_test,y_train,y_test)  
#dico_name_model
```

```

=====
RandomForestClassifier
=====
[[ 399   80]
 [   2 1914]]
precision    recall  f1-score   support

      0       0.9950      0.8330      0.9068         479
      1       0.9599      0.9990      0.9790        1916

 accuracy          0.9658         2395
 macro avg       0.9774      0.9160      0.9429         2395
weighted avg       0.9669      0.9658      0.9646         2395

=====
AdaBoost
=====
[[ 364  115]
 [  17 1899]]
precision    recall  f1-score   support

      0       0.9554      0.7599      0.8465         479
      1       0.9429      0.9911      0.9664        1916

 accuracy          0.9449         2395
 macro avg       0.9491      0.8755      0.9065         2395
weighted avg       0.9454      0.9449      0.9424         2395

```

```

=====
SVM
=====
[[ 324  155]
 [   8 1908]]
precision    recall  f1-score   support

      0       0.9759      0.6764      0.7990         479
      1       0.9249      0.9958      0.9590        1916

 accuracy          0.9319         2395
 macro avg       0.9504      0.8361      0.8790         2395
weighted avg       0.9351      0.9319      0.9270         2395

=====
KNN
=====
[[ 346  133]
 [  13 1903]]
precision    recall  f1-score   support

      0       0.9638      0.7223      0.8258         479
      1       0.9347      0.9932      0.9631        1916

 accuracy          0.9390         2395
 macro avg       0.9492      0.8578      0.8944         2395
weighted avg       0.9405      0.9390      0.9356         2395

```




```
=====
DecisionTreeClassifier
=====
[[ 407  72]
 [ 39 1877]]
```

	precision	recall	f1-score	support
0	0.9126	0.8497	0.8800	479
1	0.9631	0.9796	0.9713	1916
accuracy			0.9537	2395
macro avg	0.9378	0.9147	0.9256	2395
weighted avg	0.9530	0.9537	0.9530	2395

In order to reduce the analysis features we will use the classifier `RandomForestClassifier()` and its feature importances, we refer to the [show importances plot\(\)](#) function:

```
model_rfc = RandomForestClassifier(random_state=10)
model_rfc_fit= model_rfc.fit(X_train,y_train)
#
#-----
#--- feature_importances ---> for RandomForestClassifier
#-----
#
feature_importances_rfc=show_importances_plot(model_rfc_fit,i_threshold=0.05) #0.01
len(feature_importances_rfc)
```



We take the features that correspond to values of importance greater than the threshold 0.5 :

```
#columns list  
list(feature_importances_rfc.index
```

```
['dti', 'revol.bal', 'days.with.cr.line', 'fico',  
'inq.last.6mths']
```

```
list_rest_cols=list(feature_importances_rfc.index)  
#  
X_train1=X_train[list_rest_cols]  
X_test1=X_test[list_rest_cols]  
y_train1=y_train  
y_test1=y_test  
#  
model_rfc = RandomForestClassifier(random_state=10)  
model_rfc_fit1= model_rfc.fit(X_train1,y_train1)  
model_rfc_fit1.score(X_test1,y_test1)
```

```
0.9899791231732776
```

By restricting ourselves to these **five features**, the performance does not decrease and therefore, **for the rest we remain on this consideration:**

```
X_train1.shape,y_train1.shape,X_test1.shape,y_test1.shape
```

```
((7183, 5), (7183,), (2395, 5), (2395,))
```

We check the keras backend used :

```
# backend keras = ? tensorflow
import keras
keras.backend.backend()
```

```
tensorflow
```

We standardize the data :

```
# X_train1,X_test1,y_train1,y_test1
X_train2=X_train1.copy(deep=True)
X_test2=X_test1.copy(deep=True)
y_train2=y_train1.copy(deep=True)
y_test2=y_test1.copy(deep=True)
#
scaler2=StandardScaler()
scaler2.fit(X_train2)
#
X_train2 = scaler2.transform(X_train2)
X_test2 = scaler2.transform(X_test2)
```

The definition of the callbacks class which will be used to stop the training once the performance objective has been reached, and to display the intermediate performance after each 1000 epochs :

```
# Callback
class myCallback(tf.keras.callbacks.Callback):
    max_val_Acc=0

    def __init__(self,i_threshold):
        self.threshold_cb=i_threshold

    def on_epoch_end(self, epoch, logs={}):
        if (logs.get('val_accuracy') > self.threshold_cb):
            print("\nReached ", self.threshold_cb*100 ,
                  "% val_accuracy so cancelling training!")
            print("\nepoch: ", epoch)
            print('Acc=', logs.get('accuracy'))
            print('val_Acc=', logs.get('val_accuracy'))
            print('loss=', logs.get('loss'))
            print('\n')
            self.model.stop_training = True
        if (epoch%1000==0):
            print("\nepoch: ", epoch)
```

```
print('Acc=', logs.get('accuracy'))
print('val_Acc=', logs.get('val_accuracy'))
if logs.get('val_accuracy') > self.max_val_Acc:
    self.max_val_Acc= logs.get('val_accuracy')
print('max_val_Acc=', self.max_val_Acc)
print('loss=', logs.get('loss'))
```

for more efficiency we have defined the **InverseTimeDecay** object, and the function **get_optimizer()** to set the Adam optimizer parameters :

```
# Building model preparation
#N_VALIDATION = int(1e3)
N_TRAIN = int(1e4)
BUFFER_SIZE = int(1e4)
BATCH_SIZE = 1024
STEPS_PER_EPOCH = N_TRAIN//BATCH_SIZE

# InverseTimeDecay
lr_schedule = tf.keras.optimizers.schedules.InverseTimeDecay(
    0.001,
    decay_steps=STEPS_PER_EPOCH*300, # 1000
    decay_rate=1,
    staircase=False)

# optimizer
def get_optimizer(i_lr_schedule):
    l_adam= tf.keras.optimizers.Adam (
        learning_rate=i_lr_schedule,
        beta_1=0.9,
        name='adam'
    )
    return l_adam
```

Now, we have all the elements to launch the training with the definition of the model which is given by [get_model\(\)](#):

Model with 3 layers :

- Dense(1024, input_dim=i_input_dim,activation='relu',
kernel_regularizer=l2(i_regul_l2)))
- Dropout(0.7)
- Dense(1,activation='sigmoid'))

```
# Building model
#
target_score=0.98
callbacks = myCallback(target_score)

model_tf = get_model(X_train2.shape[1],i_def_mode=1)

model_tf.compile(loss='binary_crossentropy', optimizer=get_optimizer(lr_schedule),
                 metrics=['accuracy'])

import time
start=time.time()

history=model_tf.fit(X_train2,y_train2,epochs=50000,batch_size=BATCH_SIZE,
                    validation_data=(X_test2, y_test2) ,callbacks=[callbacks],verbose=0)

end=time.time()
print('running time: ',end-start)

score=model_tf.evaluate(X_test2,y_test2,batch_size=BATCH_SIZE)

print('\n score: ',score,'\n')
```

```
epoch: 4000
Acc= 0.9700682163238525
val_Acc= 0.9774530529975891
max_val_Acc= 0.9774530529975891
loss= 0.1011856868631058

epoch: 5000
Acc= 0.9722957015037537
val_Acc= 0.9774530529975891
max_val_Acc= 0.9774530529975891
loss= 0.09867680817842484

Reached 98.0 % val_accuracy so cancelling training!

epoch: 5939
Acc= 0.973827064037323
val_Acc= 0.9803757667541504
loss= 0.09569501876831055

running time: 1048.621102809906
3/3 [=====] - 0s 7ms/step - loss: 0.1025 - accuracy: 0.9804

score: [0.10246644914150238, 0.9803757667541504]
```

The goal for accuracy to exceed the 98% threshold was reached after 5939 epochs. We give below the confusion matrix and other classification metrics such as precision and recall. We clearly see a good result with balanced metrics.

```
print('score: ',score,'\n')
model= fit_and_evaluate_model(model_tf,X_train2,X_test2,y_train2,y_test2,i_use_fit
=False,cas_tensorflow=True)
```

```
score: [0.10246644914150238, 0.9803757667541504]

[[ 447   32]
 [  15 1901]]
      precision    recall  f1-score   support

         0       0.9675    0.9332    0.9501        479
         1       0.9834    0.9922    0.9878       1916

 accuracy                0.9804        2395
 macro avg              0.9755    0.9627    0.9689        2395
 weighted avg           0.9803    0.9804    0.9802        2395
```

We finish our presentation by displaying the evolution of the accuracy and the loss by the advancement of the epochs :



```
#-----
# Plot training and validation accuracy per epoch
#-----
acc=history.history['accuracy']
val_acc=history.history['val_accuracy']

epochs=range(len(acc)) # Get number of epochs

plt.plot(epochs, acc, 'r')
plt.plot(epochs, val_acc, 'b')
#
plt.title('Training and Validation Accuracy')
plt.legend(['Acc', 'Val_acc'])
plt.show()
print("")

#-----
# Plot training and validation loss per epoch
#-----
loss=history.history['loss']
val_loss=history.history['val_loss']
epochs=range(len(loss)) # Get number of epochs
#
plt.plot(epochs, loss, 'r')
```

```
plt.plot(epochs, val_loss, 'b')  
#  
plt.title('Training and Validation loss')  
plt.legend(['loss', 'val_loss'])  
plt.show()  
print("")
```



FUNCTIONS

Librairies

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense, Dropout
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import make_pipeline
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.metrics import confusion_matrix, classification_report
from keras.regularizers import l2
from sklearn.model_selection import train_test_split

# list classifiers
lst_algo=['AdaBoost','SVM','KNN','DecisionTreeClassifier','RandomForestClassifier']
```

test_classifierModels_list() function

```
#####
### test_classifierModels_list function
#####
def test_classifierModels_list(i_classifierModels_list,i_X_train,i_X_test,i_y_train,i_y_test):
    l_dico_name_model={}
```



```
preprocessor = make_pipeline(PolynomialFeatures(2, include_bias=False),
                             SelectKBest(f_classif, k=10))
RandomForest = make_pipeline(preprocessor, RandomForestClassifier(random_state=0))
AdaBoost = make_pipeline(preprocessor, AdaBoostClassifier(random_state=0))
SVM = make_pipeline(preprocessor, StandardScaler(), SVC(random_state=0))
KNN = make_pipeline(preprocessor, StandardScaler(), KNeighborsClassifier())
DTC = make_pipeline(preprocessor, StandardScaler(), DecisionTreeClassifier())
dict_of_models = {'RandomForestClassifier': RandomForest,
                  'AdaBoost': AdaBoost,
                  'SVM': SVM,
                  'KNN': KNN,
                  'DecisionTreeClassifier': DTC
                  }
for name, model in dict_of_models.items():
    if name in i_classifierModels_list:
        print('=====')
        print(name)
        print('=====')
        l_model= fit_and_evaluate_model(model,i_X_train,i_X_test,i_y_train,i_y_test)
        l_dico_name_model[name]=l_model
        #
return l_dico_name_model
```

fit_and_evaluate_model () function

```
#####
### Evaluation function: fit_and_evaluate_model()
#####
# Used to Evaluate models
def fit_and_evaluate_model(i_model,i_X_train,i_X_test,i_y_train,i_y_test, i_use_fit=True,cas_tensorflow=False):
    #
    l_model=i_model
    if not cas_tensorflow:
        if i_use_fit:
            l_model.fit(i_X_train, i_y_train)
        l_ypred = l_model.predict(i_X_test)
    #
```

```
if cas_tensorflow:
    l_ypred_t=np.array([0 for i in range(len(l_ypred))])
    for i in range(len(l_ypred)):
        if l_ypred[i] > 0.5:
            l_ypred_t[i]=1
    l_ypred=l_ypred_t
#
print(confusion_matrix(i_y_test, l_ypred))
print(classification_report(i_y_test, l_ypred, digits=4))

return l_model
```

get_value_counts() function

```
# list of columns with count of its values by type of columns
def get_value_counts(i_df,i_type):
    lst_nbre_val_col=[]
    cmpt=0
    for col in i_df.select_dtypes(i_type).columns:
        cmpt=i_df[col].value_counts().count()
        lst_nbre_val_col.append((col,cmpt))
#
return lst_nbre_val_col
```

show_importances_plot() function

```
#-----
#--- function show_importances_plot
#-----
# Only importances > i_threshold
def show_importances_plot(i_model, i_threshold,i_figsize=(11, 6)):
    l_feature_importances_df=pd.DataFrame({'importances': i_model.feature
_importances_,
                                           'columns_': i_model.feature_na
mes_in_})
    l_feature_importances_df.sort_values(by=['importances'], ascending=Tr
ue, inplace=True)
    feature_importances_df_t=l_feature_importances_df[l_feature_importanc
es_df['importances'] > i_threshold]
```

```
feature_importances_df_t.set_index('columns_',inplace=True)
#
feature_importances_df_t.plot(figsize=i_figsize,kind='barh')
return feature_importances_df_t
```

verify_dependency() function

```
# return list of dep columns couple
def verify_dependency(i_corr_matrix,i_threshold):
    # looking for features corresponding at the threshold 0.35
    cols=i_corr_matrix.columns
    l_tupe_dep=[] # list tuple of columns to verify dep condition
    l_cmpt=0
    for col1 in cols:
        for col2 in cols:
            l_val=i_corr_matrix.loc[col2,col1]
            #if l_cmpt < 10:
            #print (abs(l_val))
            #l_cmpt=l_cmpt+1
            if abs(l_val) > i_threshold:
                if col1!=col2:
                    if (col1,col2) not in l_tupe_dep:
                        l_tupe_dep.append((col1,col2))

    return l_tupe_dep
```

get_model() function

```
#-----
#--- function get_model()
# give model with appropriate input_dim as parameter
#-----
#
def get_model(i_input_dim,i_regul_l2=0.0001,i_def_mode=1):
    l_model=Sequential()
    if i_def_mode==1:
        l_model.add(Dense(1024, input_dim=i_input_dim,activation='relu',
            kernel_regularizer=l2(i_regul_l2)))
        l_model.add(Dropout(0.7))
    #
```

```
if i_def_mode==2:
    l_model.add(Dense(1278, input_dim=i_input_dim, activation='relu',
        kernel_regularizer=l2(i_regul_l2)))
    l_model.add(Dropout(0.5))
    l_model.add(Dense(512, activation='elu',
        kernel_regularizer=l2(i_regul_l2)),
    l_model.add(Dropout(0.5))
l_model.add(Dense(1, activation='sigmoid'))
#
```